

## Lab 6.7: Implementing a priority queue with consumer.assign()

Welcome to the session 6 lab 7. The work for this lab is done in `~/kafka-training/lab6.7`. In this lab, you are going to implement a priority queue with `consumer.assign()`.

Please refer to the [Kafka course notes](#) for any updates or changes to this lab.

Find the latest version of this lab [here](#).

### Lab Using consumer.assign to implement a priority queue

In this lab, you will implement a priority processing queue. You will use `consumer.partitionsFor(TOPIC)` to get a list of partitions. Usage like this is simplest when the partition assignment is also done manually using `assign()` instead of `subscribe()`. Use `assign()`, pass a `TopicPartition` from the consumer worker. Use the `Partitioner` from an earlier example for Producer so only important stocks get sent to the important partition.

### Using partitionsFor() for Priority Queue

`~/kafka-training/lab6.7/src/main/java/com/cloudurable/kafka/consumer/ConsumerMain.java`

Kafka Consumer: ConsumerMain.main

```
package com.cloudurable.kafka.consumer;

import com.cloudurable.kafka.StockAppConstants;
import com.cloudurable.kafka.model.StockPrice;
import org.apache.kafka.clients.consumer.Consumer;
import org.apache.kafka.clients.consumer.ConsumerConfig;
import org.apache.kafka.clients.consumer.KafkaConsumer;
import org.apache.kafka.common.serialization.StringDeserializer;
import org.slf4j.Logger;
import org.slf4j.LoggerFactory;

import java.util.ArrayList;
import java.util.List;
import java.util.Properties;
import java.util.concurrent.ExecutorService;
import java.util.concurrent.TimeUnit;
import java.util.concurrent.atomic.AtomicBoolean;
import java.util.stream.IntStream;

import static com.cloudurable.kafka.StockAppConstants.TOPIC;
import static java.util.concurrent.Executors.newFixedThreadPool;

public class ConsumerMain {
    ...
    public static void main(String... args) throws Exception {
        final AtomicBoolean stopAll = new AtomicBoolean();
        final Consumer<String, StockPrice> consumer = createConsumer();

        //Get the partitions
        final List<PartitionInfo> partitionInfos = consumer.partitionsFor(TOPIC);

        final int threadCount = partitionInfos.size();
        final int numWorkers = 5;
        final ExecutorService executorService = newFixedThreadPool(threadCount);

        IntStream.range(0, threadCount).forEach(index -> {
            final PartitionInfo partitionInfo = partitionInfos.get(index);
            final boolean leader = partitionInfo.partition() == partitionInfos.size() - 1;
            final int workerCount = leader ? numWorkers * 3 : numWorkers;
            final StockPriceConsumerRunnable stockPriceConsumer =
                new StockPriceConsumerRunnable(partitionInfo, createConsumer(),
                    readCountStatusUpdate: 10, index, stopAll, workerCount);
            consumerList.add(consumer);
            executorService.submit(stockPriceConsumer);
        });
    }
}
```

```

Runtime.getRuntime().addShutdownHook(new Thread(() -> {
    logger.info("Stopping app");
    stopAll.set(true);
    sleep();
    consumerList.forEach(Consumer::wakeup);
    executorService.shutdown();
    try {
        executorService.awaitTermination(5_000, TimeUnit.MILLISECONDS);
        if (!executorService.isShutdown())
            executorService.shutdownNow();
    } catch (InterruptedException e) {
        logger.warn("shutting down", e);
    }
    sleep();
    consumerList.forEach(Consumer::close);
}));
...
}

```

Notice that the index is the topic partition. Num threads are the partition count, and the priority partition gets extra workers.

## Using assign() for Priority Queue

~/kafka-training/lab6.7/src/main/java/com/cloudurable/kafka/consumer/StockPriceConsumerRunnable.java

**Kafka Consumer: StockPriceConsumerRunnable.runConsumer**

```

package com.cloudurable.kafka.consumer;

import com.cloudurable.kafka.model.StockPrice;
import org.apache.kafka.clients.consumer.*;
import org.apache.kafka.common.TopicPartition;
import org.slf4j.Logger;
import org.slf4j.LoggerFactory;

import java.util.Collections;
import java.util.Map;
import java.util.concurrent.*;
import java.util.concurrent.atomic.AtomicBoolean;

import static com.cloudurable.kafka.StockAppConstants.TOPIC;

public class StockPriceConsumerRunnable implements Runnable {
    ...
    private void runConsumer() throws Exception {
        //Assign a partition
        consumer.assign(Collections.singleton(topicPartition));
        final Map<String, StockPrice> lastRecordPerStock = new HashMap<>();
        try {
            int readCount = 0;
            while (isRunning()) {
                pollRecordsAndProcess(lastRecordPerStock, readCount);
            }
        } finally {
            consumer.close();
        }
    }
    ...
}

```

## Lab Work

Use the slides for Session 6 as a guide.

**ACTION - EDIT** `src/main/java/com/cloudurable/kafka/consumer/ConsumerMain.java`  
and follow the instructions in the file.

**ACTION - EDIT**

`src/main/java/com/cloudurable/kafka/consumer/StockPriceConsumerRunnable.java`  
and follow the instructions in the file.

**ACTION - RECREATE** the topic with five partitions (HINT: `bin/create-topic.sh`) and use 5 partitions.

**ACTION - RUN ZooKeeper and Brokers** if needed.

**ACTION - RUN ConsumerMain** from IDE

**ACTION - RUN StockPriceKafkaProducer** from IDE

**ACTION - OBSERVE** and then **STOP** consumers and producer

### Expected behavior

It should run and should get messages like this:

#### Expected output

```
New ConsumerRecords par count 1 count 153, max offset
ticker IBM price 66.59 Thread 4
ticker UBER price 241.94 Thread 4

New ConsumerRecords par count 1 count 220, max offset
ticker ABC price 95.85 Thread 2
ticker BBB price 53.36 Thread 2
ticker FFF price 70.34 Thread 2

New ConsumerRecords par count 1 count 318, max offset
ticker GOOG price 458.44 Thread 0
ticker DDD price 68.38 Thread 0
ticker SUN price 91.90 Thread 0
ticker INEL price 65.94 Thread 0

New ConsumerRecords par count 1 count 364, max offset
ticker AAA price 66.53 Thread 1
ticker DEF price 65.94 Thread 1
ticker EEE price 70.34 Thread 1
ticker XYZ price 65.94 Thread 1
```

### Try the following

Try using different worker pool sizes and different consumer thread pool sizes. Try adding a small wait for the processing. Try 10ms. It should all run. Stop consumer and producer when finished.

---

## Kafka Tutorial

This comprehensive *Kafka tutorial* covers Kafka architecture and design. The *Kafka tutorial* has example Java Kafka producers and Kafka consumers. The *Kafka tutorial* also covers Avro and Schema Registry.

[Complete Kafka Tutorial: Architecture, Design, DevOps and Java Examples.](#)

- [Kafka Tutorial Part 1: What is Kafka?](#)
- [Kafka Tutorial Part 2: Kafka Architecture](#)
- [Kafka Tutorial Part 3: Kafka Topic Architecture](#)
- [Kafka Tutorial Part 4: Kafka Consumer Architecture](#)
- [Kafka Tutorial Part 5: Kafka Producer Architecture](#)
- [Kafka Tutorial Part 6: Using Kafka from the command line](#)

- [Kafka Tutorial Part 7: Kafka Broker Failover and Consumer Failover](#)
  - [Kafka Tutorial Part 8: Kafka Ecosystem](#)
  - [Kafka Tutorial Part 9: Kafka Low-Level Design](#)
  - [Kafka Tutorial Part 10: Kafka Log Compaction Architecture](#)
  - [Kafka Tutorial Part 11: Writing a Kafka Producer example in Java](#)
  - [Kafka Tutorial Part 12: Writing a Kafka Consumer example in Java](#)
  - [Kafka Tutorial Part 13: Writing Advanced Kafka Producer Java examples](#)
  - [Kafka Tutorial 14: Writing Advanced Kafka Consumer Java examples](#)
  - [Kafka Tutorial Part 15: Kafka and Avro](#)
  - [Kafka Tutorial Part 16: Kafka and Schema Registry](#)
  - [Kafka Tutorial](#)
- 

#### **About Cloudurable**

We hope you enjoyed this article. Please provide [feedback](#). Cloudurable provides [Kafka training](#), [Kafka consulting](#), [Kafka support](#) and helps [setting up Kafka clusters in AWS](#).